Baseline agent: evaluation summary

Recorded Promptfoo v2 run | 2 November 2026 | 114 questions

Strict pass rate was **80.8% excluding formatting (84/104)** and **73.7% overall (84/114)**: 84 passes, 29 assertion failures and one runtime error. Strict pass requires every applicable quality, workflow, validity and formatting check. The results support a cost-reduction experiment with a smaller model and simpler workflow; savings and quality retention still need validation.

Baseline and retrieval specification

Models: gpt-5-mini for question analysis, context assessment and answering; gpt-5-mini also grades answer quality. **Workflow:** LangChain/LangGraph analyses the question, discovers company/period coverage, retrieves evidence, checks sufficiency and refines the query for up to three rounds, then answers, clarifies or refuses. LangSmith records nested runs. **Retrieval:** text-embedding-3-small (1,536 dimensions), cosine-vector search in company-specific PostgreSQL/pgvector tables; up to eight chunks per company per round, with source/page metadata. Corpus: 100 PDFs from Alphabet, Amazon, Apple, Meta and Netflix.

Answer latency: mean 33.83 s; median 31.49 s; p95 69.04 s (113 completed answers). Full run: 18 min 41 s at four-case concurrency.
Tokens: 1,611,370 agent + 137,844 judge + 4,455 embeddings = **1,753,669 total**.
Estimated API spend: $0.9642 agent including embeddings + $0.1135 grading = **$1.0777 total**, including the failed case; excludes ETL, infrastructure and taxes.

Strict pass rate by section

Section	Pass rate	Observed reason / interpretation
Factual accuracy	6/10 | 60%	Four clarifications instead of factual answers.
Tool use / coverage	13/14 | 92.9%	Google Cloud evidence was not recovered.
Out-of-domain	10/10 | 100%	Bounded refusals passed the revised judge.
Multiple companies	10/10 | 100%	Comparisons passed current checks.
Conditional reasoning	8/10 | 80%	Missing evidence / missing period discovery.
Unavailable data	7/10 | 70%	Scope errors, invalid year text, one runtime error.
Calculations	9/10 | 90%	One unnecessary fiscal/calendar clarification.
Ambiguous questions	8/10 | 80%	Clarification wording and route-contract mismatch.
Images / charts	5/10 | 50%	Wrong panel/segment, missed slides, clarification.
Table lookup	8/10 | 80%	Coverage lookup and clarification failures.
Output formatting	0/10 | 0%	Citation instructions conflict; 7/10 quality passes.

Interpretation and next steps

Main accuracy issue - over-clarification. Requests such as "earnings in 2025" trigger revenue-versus-net-income and fiscal-versus-calendar questions. Many failures are failure to answer, rather than wrong arithmetic. Define metric/period defaults, state assumptions, and clarify when uncertainty prevents a defensible answer. The judge accepted some incomplete clarifications, so quality-only scores need manual review.

Formatting - prompt/evaluation conflict. The system mandates citations while users request a bare number or company name. Proposed change: honor explicit user formats and retain provenance separately. Assess the resulting instruction-following and prompt-injection risks in the next-generation agent; no prompt change was made for this report.

Data quality versus agent error remains unresolved. Image accuracy is weak; table lookup is 80%, with observed failures mainly in coverage/clarification. Inspect original PDF -> extracted table/figure description -> chunk -> retrieved context -> answer to isolate ETL loss, retrieval misses or generation mistakes. Apply the same audit to factual misses before attributing them to the model or scaling it down.

Source: saved v2 summary, cases and run manifest in reports/baseline-2026-11-02-v2/. Recorded scores; AMB04 was reclassified afterward. No tests/evaluations rerun. One trial per case; judge errors remain.